

USE CASE STUDY REPORT

Group No.: 10

Student Names: Srishti Bhandari and Fangsheng Yan

Data Mining techniques to predict IMDB scores

Executive Summary:

Movies have become a favorite way of entertainment over the centuries. There are thousands of movies coming out each year. IMDB score is a popular and important indicator for people to determine the success of movies. However, it is not certain that movies with high budget or gross earnings will have higher IMDB scores than a fairly lower budget movie with new actors will have a lower IMDB score.

Goal of the study:

The success of a movie isn't purely about entertainment of the audience, the producing firms, directors, actors and the entire crew make huge profits from them. It is essential to keep in mind the kind of profits the movie will make before investing money in the making. Our project aims at providing insights on whether the movie will gain a high IMDB score based on multitude of factors that contribute towards this.

Origin of the data and data processing:

The data is obtained from a verified website: Kaggle. The data consists of 28 variables for 5043 movies, spanning across 100 years in 66 countries. There are about 2399 unique director names and thousands of actors and actresses. The "imdb_score" variable is a column in our data which is the response variable while the other 27 variables are the possible predictors. The data does contain missing records and special characters which gives us an opportunity to apply principals of data cleansing and pre-processing. Various Packages and functions are used for the same.

Data Mining techniques:

The project shows the movie ratings in 3 parameters: Bad, ok, good and excellent. To figure out whether a movie will do good or bad, it is essential to run the models through algorithms of data mining. We have used: Classification tree, K Nearest neighbor and Random forest to process our models. The data for KNN and Random Forest has been divided into 3: training, validation and test.

I. Background and Introduction

With the Boom of technology, the art of movie making has improved by folds. The story telling has seen changes as well. Movies are a source of entertainment which provides entertainment to everyone irrespective of age and culture. The reason for choosing this topic as our project was to learn algorithms using a dataset that we everyone could relate too and use it with real world examples of things/movies we have watched.

The problem we are trying to resolve is the uncertainty of the movie performing well or not. It is imperative to have some idea of whether the movie is going to make a hit or be a flop or perform somewhere in the middle.

IMDB website has proven to be a very reliable source for providing information to the users based on multiple factors and reviews from audience as well as critics.


The goal is to identify a pattern of factors which contribute to the IMDB score the movie receives. The factors are: directors, actors, location, language the movie is in, the country it is produced etc.

We aim to firstly clean this data and put it into a format which can be easily understood and used to make modifications on. We then use data visualization techniques to further analyze the data at hand and understand the relation between factors. Moving forward we apply three data mining models to determine the factors contributing to their respective IMDB scores and the accuracy of each of them to find the best suited one.

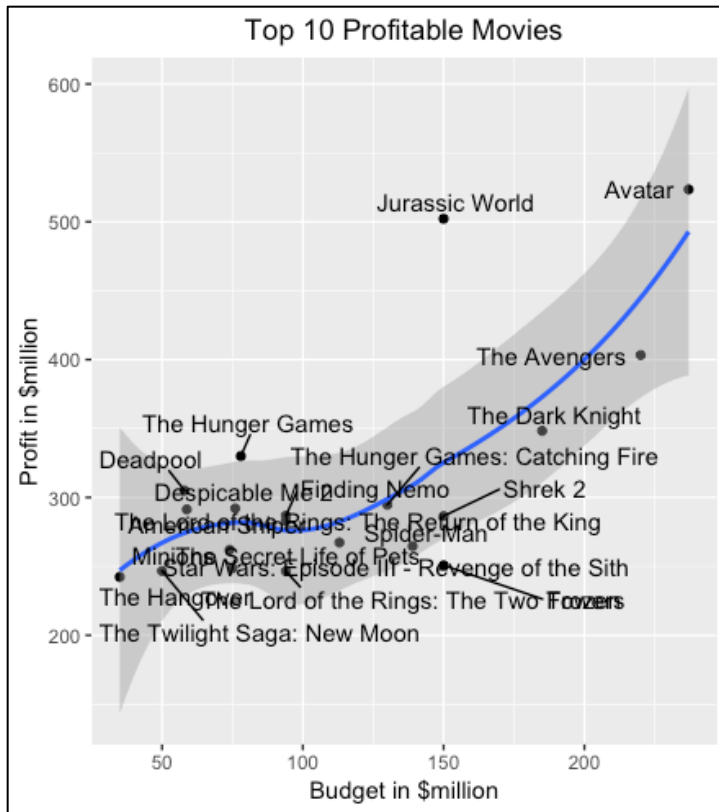
II. Data Exploration and Visualization

Provide brief description of techniques used to explore the data including: basic charts, distribution plots, correlations, missing values, rescaling, aggregation, hierarchies, zooming, filtering, etc.

1. Plotting histogram of movies released over the years

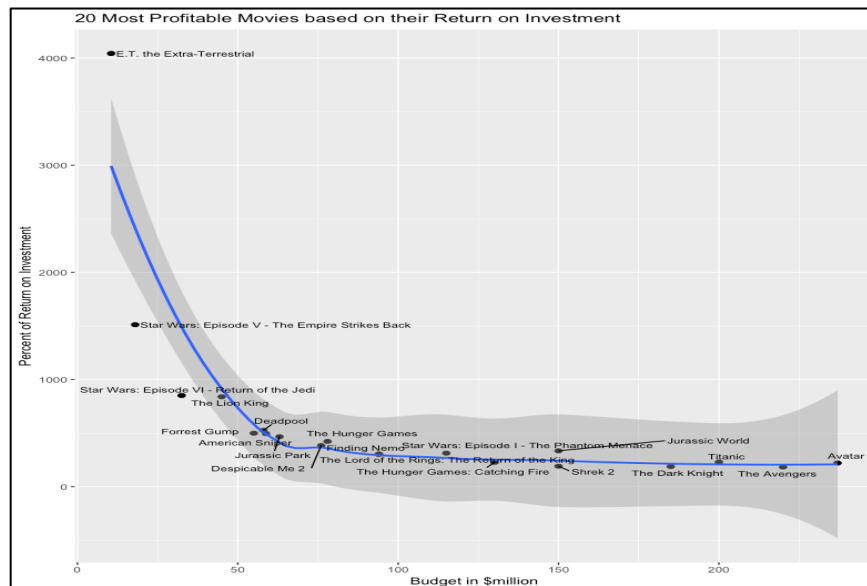


2. It can be seen that the movie boom came after 1980 and thus we represent the data only after 1980.
3. We create 2 new variables: profit and return on investment (%)
 - Where profit = gross - budget,
 - And $\text{return_on_investment_perc} = (\text{profit}/\text{budget}) * 100$



Adjacent is the visualization of the top 20 movies based on profits in millions\$


4. Using profits and return on investment variables as criteria to find 20 most profitable movies



5. Visualizing 20 top directors based on the highest IMDB scores


director_name	avg_imdb
Tony Kaye	8.600000
Damien Chazelle	8.500000
Majid Majidi	8.500000
Ron Fricke	8.500000
Christopher Nolan	8.425000
Asghar Farhadi	8.400000
Marius A. Markevicius	8.400000
Richard Marquand	8.400000
Sergio Leone	8.400000
Lee Unkrich	8.300000
Lenny Abrahamson	8.300000
Pete Docter	8.233333
Hayao Miyazaki	8.225000
Joshua Oppenheimer	8.200000
Juan José Campanella	8.200000
Quentin Tarantino	8.200000
David Singleton	8.100000
Je-kyu Kang	8.100000
Terry George	8.100000
Tim Miller	8.100000

6. Plotting a visualization for commercial success vs critical claim



The above observation shows that there is hardly any correlation between critical acclaim and the movie's commercial success

7. Visualizing relation between Facebook likes and IMDB scores



Movies with high Facebook likes can be seen to have higher IMDB score.

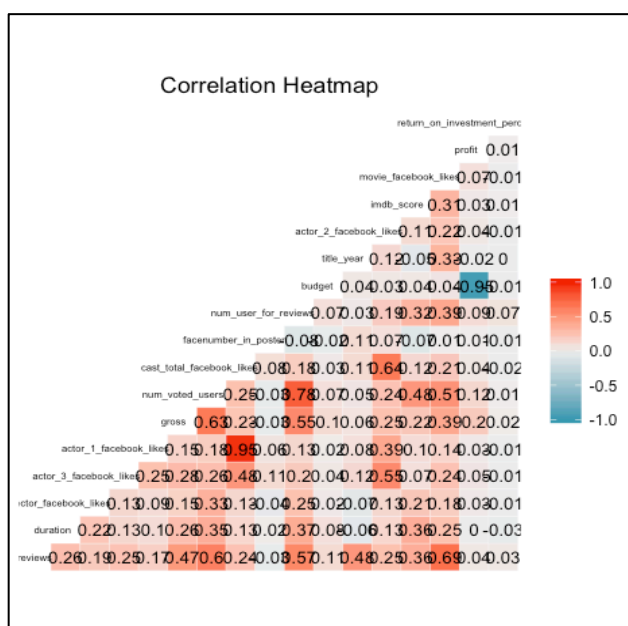
Exploration of Data:

- Find the total number of directors and actors:

```
> sum(uniqueN(movie_data$director_name))
[1] 1660
> sum(uniqueN(movie_data[, c("actor_1_name", "actor_2_name", "actor_3_name")]))
[1] 3621
```

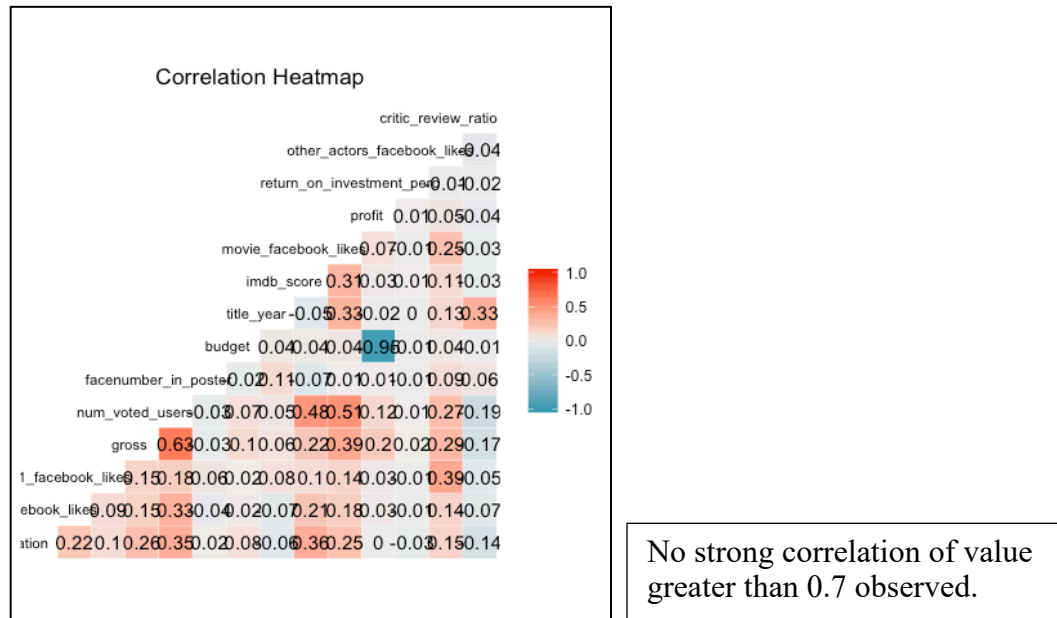
The names of the actors and directors is so distinct that they can't possibly contribute towards predicting.

- The movie IMDB link is redundant.
- To avoid multicollinearity, we remove the 2 previously added variables
- Plotting the map for the remaining values



- Based on the heatmap, we can see some high correlations (>0.7) between predictors. The highest correlation value observed is 0.95 and we can see that actor_1_facebook_likes is highly correlated with the cast_total_facebook_likes and both actor2 and actor3 are also correlated to the total. Thus we modify them into 2 variables: actor_1_facebook_likes and other_actors_facebook_likes.

5. Plotting the heatmap again post deleting the data



- The aim is to build a project wherein the model predicts whether the movie is good or bad. So, bin the scores in four buckets: less than 4(Bad), 4-6(OK), 6-8(Good) and 8-10(Excellent)


III. Data Preparation and Preprocessing

We use various methods for data processing and pre-processing, to remove unnecessary data / data with special characters.

We follow the process:

- Look for duplicates and delete them – remake the data frame by adding only unique values.
- Movie title across the data has a special character which needs to be corrected- remove the special character and remake the data frame.
- The genres are widespread and for better understanding we split the genres into multiple data frames. We find the IMDB scores for each of the genres and plot the means in a Bar-plot.
- We see that the means are all in the range 6-8 and thus it can be concluded that genres is not an impacting factor towards the IMDB score.

5. We then find the aggregate of NAs in the data – plot a heatmap to visualize all the missing values.



```
> colSums(sapply(movie_data, is.na))
```

color	director_name	num_critic_for_reviews	duration
0	0	1	1
director_facebook_likes	actor_3_facebook_likes	actor_2_name	actor_1_facebook_likes
0	10	0	3
gross	actor_1_name	movie_title	num_voted_users
0	0	0	0
cast_total_facebook_likes	actor_3_name	facenumber_in_poster	plot_keywords
0	0	6	0
movie_imdb_link	num_user_for_reviews	language	country
0	0	0	0
content_rating	budget	title_year	actor_2_facebook_likes
0	0	0	5
imdb_score	aspect_ratio	movie_facebook_likes	
0	74	0	

The number of NAs after removing data.

6. We observe that aspect_ratio still consists of NAs and thus we replace all NAs with 0.
 7. Find the aspect_ratio, it is observed the 1.85 and 2.35 are the most common ones and thus we find the mean of values with the same aspect ratio.

```
> mean(movie_data$imdb_score[movie_data$aspect_ratio == 1.85])
[1] 6.373938
>
> mean(movie_data$imdb_score[movie_data$aspect_ratio == 2.35])
[1] 6.508471
```

8. The mean where the aspect_ratio isn't 1.85 and 2.35

```
> mean(movie_data$imdb_score[movie_data$aspect_ratio != 1.85 & movie_data$aspect_ratio != 2.35])
[1] 6.672519
```

9. The mean in either of the cases isn't deviating much and it can be assumed that removing this variable will not affect our analysis.
10. For the variable faceumber_in_poster, we replace the NAs with the column mean and the 0s in the column are replaced with NAs
11. Convert all the 0s in the data to NA
12. Replacing NA in num_critic_for_reviews, duration, director_facebook_likes, actor_3_facebook_likes, actor_1_facebook_likes, cast_total_facebook_likes, actor_2_facebook_likes, movie_facebook_likes with the average of the column
13. Blanks are to be considered as missing values, removing all of them.
14. Observing the content rating :

```
> table(movie_data$content_rating)
```

Approved	G	GP	M	NC-17	Not Rated	Passed	PG	PG-13	R	TV-14
51	17	91	1	2	6	42	3	573	1314	1723
TV-G	TV-MA	TV-PG	TV-Y	TV-Y7	Unrated	X				
0	0	0	0	0	24	10				

15. On evaluating content ratings, we observe M = GP = PG, X = NC-17. Replace M and GP with PG and replace X with NC-17
16. Replace "Approved", "Not Rated", "Passed", "Unrated" with the most common rating "R"

```
> table(movie_data$content_rating)
```

G	NC-17	PG	PG-13	R
91	16	576	1314	1809

The new data for content rating.

17. Evaluating the color of the movies:

Black and White	Color
2	124
	3680

It can be observed that the data in color is completely partial towards colored movies and thus it is not an influential factor and we can remove it.

18. Observing the language data


```
> table(movie_data$language)
```

Aboriginal	Arabic	Aramaic	Bosnian	Cantonese	Chinese	Czech	Danish	Dari	Dutch	
2	2	1	1	7	0	1	3	2	3	
Dzongkha	English	Filipino	French	German	Greek	Hebrew	Hindi	Hungarian	Icelandic	Indonesian
0	3644	1	34	11	0	2	5	1	0	2
Italian	Japanese	Kannada	Kazakh	Korean	Mandarin	Maya	Mongolian	None	Norwegian	Panjabi
7	10	0	1	5	14	1	1	1	4	0
Persian	Polish	Portuguese	Romanian	Russian	Slovenian	Spanish	Swahili	Swedish	Tamil	Telugu
3	0	5	1	1	0	24	0	0	0	0
Thai	Urdu	Vietnamese	Zulu							
3	0	1	1							

It can be observed that the data in languages is completely partial towards English movies and thus it is not an influential factor and we can remove it.

19. Checking if the country the movie is produced in is an influential factor towards it's score

	Afghanistan	Argentina	Aruba	Australia	Bahamas
0	1	3	1	40	0
Belgium	Brazil	Bulgaria	Cambodia	Cameroon	Canada
1	5	0	0	0	63
Chile	China	Colombia	Czech Republic	Denmark	Dominican Republic
1	13	1	3	9	0
Egypt	Finland	France	Georgia	Germany	Greece
0	1	103	1	79	1
Hong Kong	Hungary	Iceland	India	Indonesia	Iran
13	2	1	5	1	4
Ireland	Israel	Italy	Japan	Kenya	Kyrgyzstan
7	2	11	15	0	0
Libya	Mexico	Netherlands	New Line	New Zealand	Nigeria
0	10	3	1	11	0
Norway	Official site	Pakistan	Panama	Peru	Philippines
4	1	0	0	1	1
Poland	Romania	Russia	Slovakia	Slovenia	South Africa
1	2	3	0	0	3
South Korea	Soviet Union	Spain	Sweden	Switzerland	Taiwan
8	0	22	0	0	2
Thailand	Turkey	UK United Arab Emirates		USA	West Germany
4	0	316	0	3025	1

Approximately 79% movies are form the US, 8% from UK and 13% from other countries. Thus, we collectively represent the movie locations as: US, UK, Others.

```
> table(movie_data$country)
```

UK	USA	Others
316	3025	465

```
> |
```

20. To avoid multicollinearity, we remove the 2 previously added variables.

IV. Data Mining Techniques and Implementation

&

V. Performance Evaluation

To apply models, splitting the data into training, validation and test sets with the ratio of 6:2:2

CLASSIFICATION TREE

Plotting a full-grown tree:

```
Classification tree:
rpart(formula = binned_score ~ . - imdb_score, data = train,
      method = "class", cp = 1e-05, minsplit = 5, xval = 5)

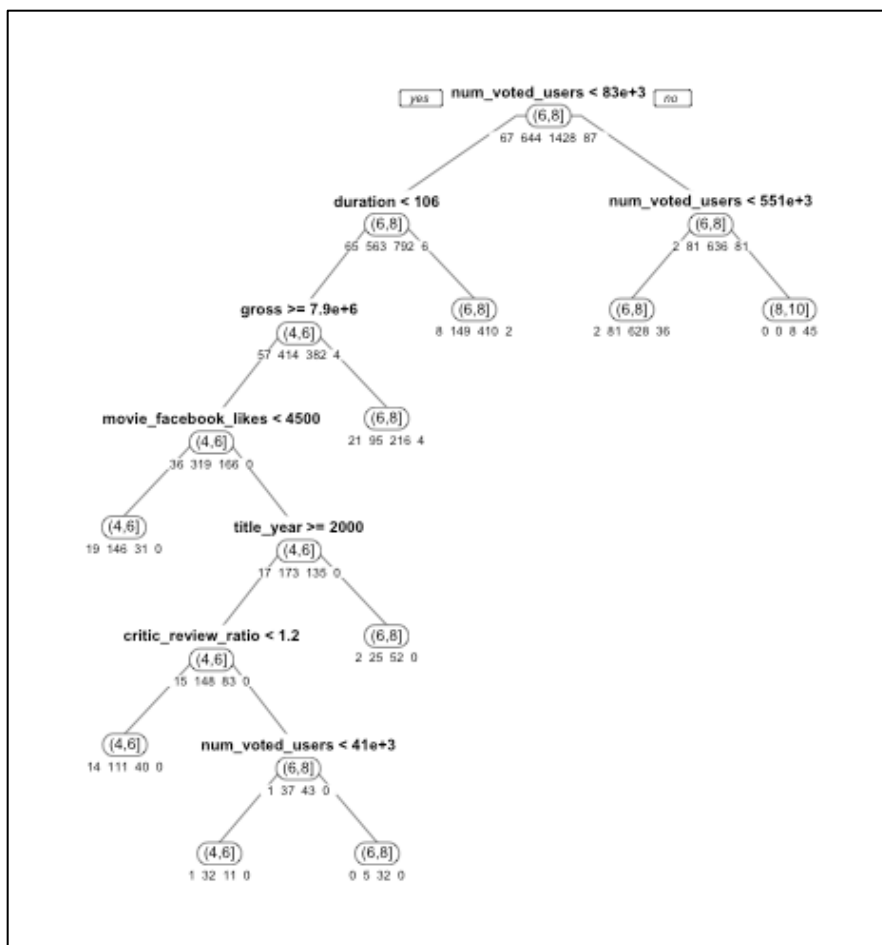
Variables actually used in tree construction:
[1] actor_1_facebook_likes    budget                content_rating        country
[5] critic_review_ratio       director_facebook_likes duration              facenumber_in_poster
[9] gross                    movie_facebook_likes  num_voted_users      other_actors_facebook_likes
[13] profit                   return_on_investment_perc title_year

Root node error: 798/2226 = 0.35849
```

```
Root node error: 798/2226 = 0.35849

n= 2226

      CP nsplit rel error  xerror   xstd
1  0.06390977      0  1.00000  1.00000  0.028353
2  0.04636591      3  0.80827  0.86216  0.027322
3  0.01691729      4  0.76190  0.79574  0.026697
4  0.00751880      8  0.69424  0.77694  0.026503
5  0.00626566     10  0.67920  0.77318  0.026464
6  0.00563910     13  0.66040  0.75689  0.026289
7  0.00543024     15  0.64912  0.75689  0.026289
8  0.00501253     20  0.61278  0.75439  0.026262
9  0.00407268     25  0.58772  0.76817  0.026411
10 0.00375940     29  0.57143  0.78195  0.026556
11 0.00325815     45  0.50877  0.79198  0.026659
12 0.00313283     50  0.49248  0.79198  0.026659
13 0.00292398     52  0.48622  0.79699  0.026709
14 0.00250627     55  0.47744  0.79825  0.026722
15 0.00187970    104  0.34461  0.80702  0.026809
16 0.00167084    131  0.27945  0.80326  0.026772
17 0.00156642    143  0.25940  0.80326  0.026772
18 0.00125313    152  0.24436  0.82957  0.027026
19 0.00093985    203  0.17794  0.83208  0.027049
20 0.00083542    207  0.17419  0.82957  0.027026
21 0.00075188    210  0.17168  0.82957  0.027026
22 0.00062657    215  0.16792  0.84586  0.027176
23 0.00027847    221  0.16416  0.84962  0.027210
24 0.00020886    230  0.16165  0.85464  0.027255
25 0.0001000     236  0.16040  0.85589  0.027266
```

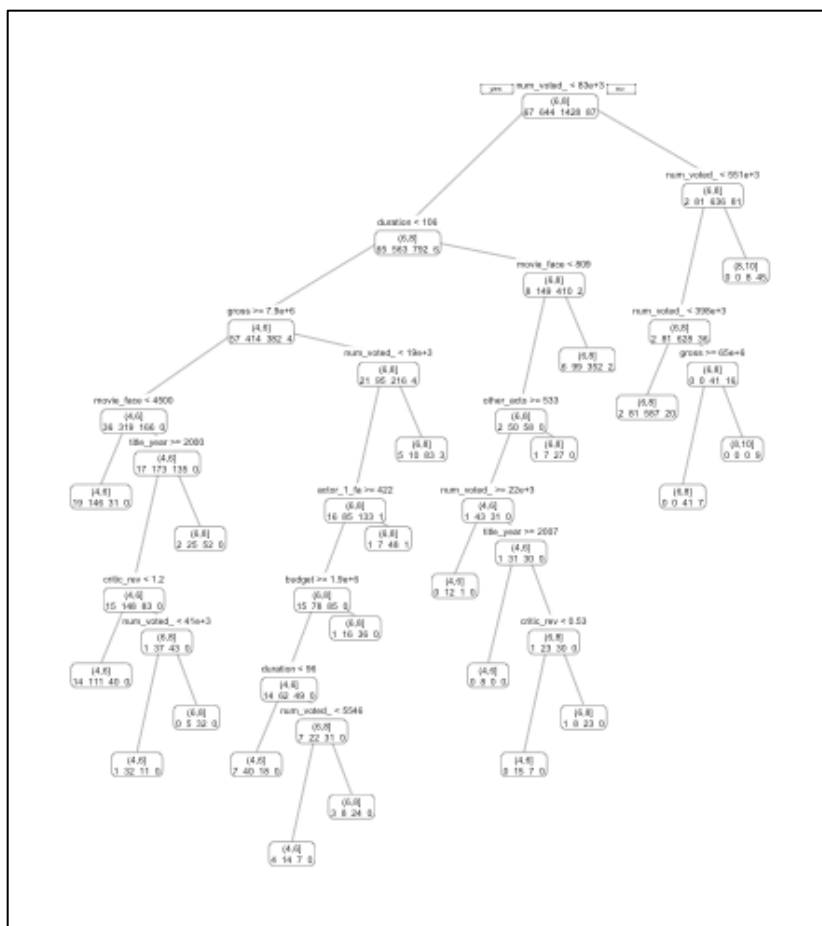


We plot pruned trees as well:

The length of the pruned tree is 21

On Test Data:

Confusion Matrix and Statistics				
	Reference			
Prediction	(0,4]	(4,6]	(6,8]	(8,10]
(0,4]	0	0	0	0
(4,6]	45	378	115	0
(6,8]	22	266	1305	33
(8,10]	0	0	8	54
Overall Statistics				
Accuracy : 0.7803				
95% CI : (0.7625, 0.7974)				
No Information Rate : 0.6415				
P-Value [Acc > NIR] : < 2.2e-16				
Kappa : 0.5228				
McNemar's Test P-Value : NA				
Statistics by Class:				
	Class: (0,4]	Class: (4,6]	Class: (6,8]	Class: (8,10]
Sensitivity	0.0000	0.5870	0.9139	0.62069
Specificity	1.0000	0.8989	0.5977	0.99626
Pos Pred Value	NaN	0.7026	0.8026	0.87097
Neg Pred Value	0.9699	0.8424	0.7950	0.98475
Prevalence	0.0301	0.2893	0.6415	0.03908
Detection Rate	0.0000	0.1698	0.5863	0.02426
Detection Prevalence	0.0000	0.2417	0.7305	0.02785
Balanced Accuracy	0.5000	0.7429	0.7558	0.80847



On Validation set:

Confusion Matrix and Statistics				
Reference				
Prediction	(0,4]	(4,6]	(6,8]	(8,10]
(0,4]	0	0	0	0
(4,6]	9	88	62	0
(6,8]	6	121	424	10
(8,10]	0	0	5	17
Overall Statistics				
Accuracy : 0.7129				
95% CI : (0.6789, 0.7453)				
No Information Rate : 0.6617				
P-Value [Acc > NIR] : 0.001616				
Kappa : 0.345				
McNemar's Test P-Value : NA				
Statistics by Class:				
	Class: (0,4]	Class: (4,6]	Class: (6,8]	Class: (8,10]
Sensitivity	0.00000	0.4211	0.8635	0.62963
Specificity	1.00000	0.8668	0.4542	0.99301
Pos Pred Value	NaN	0.5535	0.7558	0.77273
Neg Pred Value	0.97978	0.7925	0.6298	0.98611
Prevalence	0.02022	0.2817	0.6617	0.03639
Detection Rate	0.00000	0.1186	0.5714	0.02291
Detection Prevalence	0.00000	0.2143	0.7561	0.02965
Balanced Accuracy	0.50000	0.6439	0.6589	0.81132

On Test data:

Confusion Matrix and Statistics				
Reference				
Prediction	(0,4]	(4,6]	(6,8]	(8,10]
(0,4]	0	0	0	0
(4,6]	8	107	76	0
(6,8]	5	105	423	10
(8,10]	0	0	1	8
Overall Statistics				
Accuracy : 0.7241				
95% CI : (0.6904, 0.756)				
No Information Rate : 0.6729				
P-Value [Acc > NIR] : 0.001485				
Kappa : 0.3651				
McNemar's Test P-Value : NA				
Statistics by Class:				
	Class: (0,4]	Class: (4,6]	Class: (6,8]	Class: (8,10]
Sensitivity	0.0000	0.5047	0.8460	0.44444
Specificity	1.0000	0.8418	0.5062	0.99862
Pos Pred Value	NaN	0.5602	0.7790	0.88889
Neg Pred Value	0.9825	0.8098	0.6150	0.98638
Prevalence	0.0175	0.2853	0.6729	0.02423
Detection Rate	0.0000	0.1440	0.5693	0.01077
Detection Prevalence	0.0000	0.2571	0.7308	0.01211
Balanced Accuracy	0.5000	0.6733	0.6761	0.72153

K NEAREST NEIGHBOR

Training data accuracy

Validation data accuracy

```
> accuracy.df
  k accuracy
1  1 0.6671159
2  2 0.6293801
3  3 0.6981132
4  4 0.6940701
5  5 0.6994609
6  6 0.6900270
7  7 0.6886792
8  8 0.6913747
9  9 0.7008086
10 10 0.7048518
11 11 0.6954178
12 12 0.6967655
13 13 0.6954178
14 14 0.6765499
15 15 0.6994609
16 16 0.6927224
17 17 0.6981132
18 18 0.6886792
19 19 0.6927224
20 20 0.6954178
```


```
> accuracy.df
  k accuracy
1  1 0.66711
2  2 0.64555
3  3 0.70080
4  4 0.68867
5  5 0.70080
6  6 0.70350
7  7 0.69002
8  8 0.70485
9  9 0.70754
10 10 0.70485
11 11 0.69676
12 12 0.69002
13 13 0.69407
14 14 0.68328
15 15 0.69676
16 16 0.69272
```

Accuracy of confusion Matrix:


Accuracy
0.7456258

RANDOM FOREST

Visualizing the model error



Visualizing the relative importance of variables



Finding the accuracy of the model

Confusion Matrix and Statistics				
	Reference			
Prediction (0,4] (4,6] (6,8] (8,10]				
(0,4]	0	0	0	0
(4,6]	10	114	41	0
(6,8]	5	95	449	12
(8,10]	0	0	1	15
Overall Statistics				
Accuracy : 0.779				
95% CI : (0.7474, 0.8083)				
No Information Rate : 0.6617				
P-Value [Acc > NIR] : 1.804e-12				
Kappa : 0.4934				
McNemar's Test P-Value : NA				
Statistics by Class:				
	Class: (0,4]	Class: (4,6]	Class: (6,8]	Class: (8,10]
Sensitivity	0.00000	0.5455	0.9145	0.55556
Specificity	1.00000	0.9043	0.5538	0.99860
Pos Pred Value	NaN	0.6909	0.8004	0.93750
Neg Pred Value	0.97978	0.8354	0.7680	0.98347
Prevalence	0.02022	0.2817	0.6617	0.03639
Detection Rate	0.00000	0.1536	0.6051	0.02022
Detection Prevalence	0.00000	0.2224	0.7561	0.02156
Balanced Accuracy	0.50000	0.7249	0.7341	0.77708

VI. Discussion and Recommendation

Observing the accuracy, we can see that:

Dataset	Decision Tree	KNN	Random Forest
Training	0.7803		
Validation	0.7126	0.7143	0.7654
Test	0.7241	0.7456	0.779

We thus decide random forest is the best algorithm to use in this case.

VII. Summary

The dataset left after pre-processing and removing all the unnecessary data, we find the accuracy of a classification model using classification tree/pruned tree, K Nearest Neighbor and Random Forest. We found Random forest to give the highest accuracy and thus we choose this as the appropriate model.

Appendix: R Code for use case study

```
#Load Packages
install.packages("ggplot2")
library(ggplot2)
install.packages("ggrepel")
library(ggrepel)
install.packages("ggthemes")
library(ggthemes)
install.packages("scales")
```

```

library(scales)
install.packages("dplyr")
library(dplyr)
install.packages("VIM")
library(VIM)
install.packages("data.table")
library(data.table)
install.packages("formattable")
library(formattable)
install.packages("plotly")
library(plotly)
install.packages("corrplot")
library(corrplot)
install.packages("GGally")
library(GGally)
install.packages("caret")
library(caret)
install.packages("car")
library(car)

#Read Data

movie_data <- read.csv("~/Desktop/movie_metadata.csv", header=TRUE)
str(movie_data)

# DATA EXPLORATION

# Look for duplicates and delete them
sum(duplicated(movie_data))
movie_data <- movie_data[!duplicated(movie_data),]

# Today the movie title- garbage found before the actual name

library(stringr)
movie_data$movie_title <- gsub("Â", "", as.character(factor(movie_data$movie_title)))
str_trim(movie_data$movie_title, side = "right")

# Check all genres of the movies

head(movie_data$genres)

# Create a dataframe to store the substrings
genres.df <- as.data.frame(movie_data[,c("genres", "imdb_score")])

# Separate different genres

```



```

genres.df$Action <- sapply(1:length(genres.df$genres), function(x) if (genres.df[x,1]
  %like% "Action") 1 else 0)
genres.df$Action <- sapply(1:length(genres.df$genres), function(x) if (genres.df[x,1]
  %like% "Action") 1 else 0)
genres.df$Adventure <- sapply(1:length(genres.df$genres), function(x) if (genres.df[x,1]
  %like% "Adventure") 1 else 0)
genres.df$Animation <- sapply(1:length(genres.df$genres), function(x) if (genres.df[x,1]
  %like% "Animation") 1 else 0)
genres.df$Biography <- sapply(1:length(genres.df$genres), function(x) if (genres.df[x,1]
  %like% "Biography") 1 else 0)
genres.df$Comedy <- sapply(1:length(genres.df$genres), function(x) if (genres.df[x,1]
  %like% "Comedy") 1 else 0)
genres.df$Crime <- sapply(1:length(genres.df$genres), function(x) if (genres.df[x,1]
  %like% "Crime") 1 else 0)
genres.df$Documentary <- sapply(1:length(genres.df$genres), function(x) if
  (genres.df[x,1] %like% "Documentary") 1 else 0)
genres.df$Drama <- sapply(1:length(genres.df$genres), function(x) if (genres.df[x,1]
  %like% "Drama") 1 else 0)
genres.df$Family <- sapply(1:length(genres.df$genres), function(x) if (genres.df[x,1]
  %like% "Family") 1 else 0)
genres.df$Fantasy <- sapply(1:length(genres.df$genres), function(x) if (genres.df[x,1]
  %like% "Fantasy") 1 else 0)
genres.df$`Film-Noir` <- sapply(1:length(genres.df$genres), function(x) if
  (genres.df[x,1] %like% "Film-Noir") 1 else 0)
genres.df$History <- sapply(1:length(genres.df$genres), function(x) if (genres.df[x,1]
  %like% "History") 1 else 0)
genres.df$Horror <- sapply(1:length(genres.df$genres), function(x) if (genres.df[x,1]
  %like% "Horror") 1 else 0)
genres.df$Musical <- sapply(1:length(genres.df$genres), function(x) if (genres.df[x,1]
  %like% "Musical") 1 else 0)
genres.df$Mystery <- sapply(1:length(genres.df$genres), function(x) if (genres.df[x,1]
  %like% "Mystery") 1 else 0)
genres.df$News <- sapply(1:length(genres.df$genres), function(x) if (genres.df[x,1]
  %like% "News") 1 else 0)
genres.df$Romance <- sapply(1:length(genres.df$genres), function(x) if (genres.df[x,1]
  %like% "Romance") 1 else 0)
genres.df$`Sci-Fi` <- sapply(1:length(genres.df$genres), function(x) if (genres.df[x,1]
  %like% "Sci-Fi") 1 else 0)
genres.df$Short <- sapply(1:length(genres.df$genres), function(x) if (genres.df[x,1]
  %like% "Short") 1 else 0)
genres.df$Sport <- sapply(1:length(genres.df$genres), function(x) if (genres.df[x,1]
  %like% "Sport") 1 else 0)
genres.df$Thriller <- sapply(1:length(genres.df$genres), function(x) if (genres.df[x,1]
  %like% "Thriller") 1 else 0)
genres.df$War <- sapply(1:length(genres.df$genres), function(x) if (genres.df[x,1]
  %like% "War") 1 else 0)
genres.df$Western <- sapply(1:length(genres.df$genres), function(x) if (genres.df[x,1]
  %like% "Western") 1 else 0)

```

```

# Find the mean of imdb score for different genres

means <- rep(0,23)
for (i in 1:23) {
  means[i] <- mean(genres.df$imdb_score[genres.df[i+2]==1])
}

# Plot the Means

barplot(means, main = "Mean od the imdb scores for different genres")

# All means are in the range of 6-8, it can be assumed that not a lot of difference will be
  made to the
#IMDB score if genres were present

movie_data <- subset(movie_data, select = -c(genres))

# Making sure Genres returns a NULL
str(movie_data$genres)

# Find Aggregate of NAs in all columns

colSums(sapply(movie_data, is.na))

# Plotting a heat map to visualize the missing values

missing.values <- aggr(movie_data, sortVars = T, prop = T, sortCombs = T, cex.lab =
  1.5, cex.axis = .6, cex.numbers = 5, combined = F, gap = -.2)

# Gross and Budget have the highest amount of missing values but both of them are
  important factors in determing
# the IMDB score and thus we remove only the rows which have NA in them

movie_data <- movie_data[!is.na(movie_data$gross),]
movie_data <- movie_data[!is.na(movie_data$budget),]

# Checking how much was our daat affected due to removing rows

dim(movie_data)

# 23% of data removed, still consists of 3857 records for analysis
# Rechecking the number of NAs

sum(complete.cases(movie_data))

colSums(sapply(movie_data, is.na))

```

```

# aspect ratio has the highest number of NAs, checking how important aspect ration in
prediction is

table(movie_data$aspect_ratio)

# Replacing NAs in aspect ration with 0

movie_data$aspect_ratio[is.na (movie_data$aspect_ratio)] <- 0

# Checking mean where aspect ratio is 1.85 and 2.35
mean(movie_data$imdb_score[movie_data$aspect_ratio == 1.85])

mean(movie_data$imdb_score[movie_data$aspect_ratio == 2.35])

# Checking mean where aspect ratio is not 1.85 and 2.35
mean(movie_data$imdb_score[movie_data$aspect_ratio != 1.85 &
movie_data$aspect_ratio != 2.35])

# Observed: The mean in either of the cases isn't deviating much and it can be assumed
tht removing this variable
# will not affect our analysis

movie_data <- subset(movie_data, select = -c(aspect_ratio))

# Rechecking if the aspect ratio is still present or is NULL

str(movie_data$aspect_ratio)

# Replacing NAs and 0s in the Data

# Replacing NA in facenumber_in_poster with the average of the column

movie_data$facenumber_in_poster[is.na(movie_data$facenumber_in_poster)] <-
round(mean(movie_data$facenumber_in_poster, na.rm = TRUE))

# Convert 0s in the data to NAs
movie_data[,c(5,6,8,13,24,26)][movie_data[,c(5,6,8,13,24,26)] == 0] <- NA

# Replacing NA in num_critic_for_reviews with the average of the column

movie_data$num_critic_for_reviews[is.na(movie_data$num_critic_for_reviews)] <-
round(mean(movie_data$num_critic_for_reviews, na.rm = TRUE))

# Replacing NA in duration with the average of the column

movie_data$duration[is.na(movie_data$duration)] <- round(mean(movie_data$duration,
na.rm = TRUE))

```

```

# Replacing NA in director_facebook_likes with the average of the column

movie_data$director_facebook_likes[is.na(movie_data$director_facebook_likes)] <-
  round(mean(movie_data$director_facebook_likes, na.rm = TRUE))

# Replacing NA in actor_3_facebook_likes with the average of the column

movie_data$actor_3_facebook_likes[is.na(movie_data$actor_3_facebook_likes)] <-
  round(mean(movie_data$actor_3_facebook_likes, na.rm = TRUE))

# Replacing NA in actor_1_facebook_likes with the average of the column

movie_data$actor_1_facebook_likes[is.na(movie_data$actor_1_facebook_likes)] <-
  round(mean(movie_data$actor_1_facebook_likes, na.rm = TRUE))

# Replacing NA in cast_total_facebook_likes with the average of the column

movie_data$cast_total_facebook_likes[is.na(movie_data$cast_total_facebook_likes)] <-
  round(mean(movie_data$cast_total_facebook_likes, na.rm = TRUE))

# Replacing NA in actor_2_facebook_likes with the average of the column

movie_data$actor_2_facebook_likes[is.na(movie_data$actor_2_facebook_likes)] <-
  round(mean(movie_data$actor_2_facebook_likes, na.rm = TRUE))

# Replacing NA in movie_facebook_likes with the average of the column

movie_data$movie_facebook_likes[is.na(movie_data$movie_facebook_likes)] <-
  round(mean(movie_data$movie_facebook_likes, na.rm = TRUE))

# Finding the missing values in content rating

table(movie_data$content_rating)

# Blanks are to be considered as missing values

movie_data <- movie_data[!(movie_data$content_rating %in% ""),]

# Re-evaluating content ratings
# M = GP = PG, X = NC-17. Replace M and GP with PG and replace X with NC-17

movie_data$content_rating
movie_data$content_rating[movie_data$content_rating == 'M'] <- 'PG'
movie_data$content_rating[movie_data$content_rating == 'GP'] <- 'PG'
movie_data$content_rating[movie_data$content_rating == 'X'] <- 'NC-17'
# Replace "Approved", "Not Rated", "Passed", "Unrated" with the most common rating
"R"

```

```

movie_data$content_rating[movie_data$content_rating == 'Approved'] <- 'R'
movie_data$content_rating[movie_data$content_rating == 'Not Rated'] <- 'R'
movie_data$content_rating[movie_data$content_rating == 'Passed'] <- 'R'
movie_data$content_rating[movie_data$content_rating == 'Unrated'] <- 'R'
movie_data$content_rating <- factor(movie_data$content_rating)

table(movie_data$content_rating)

# Creating 2 columns profit and percentage of return based on gross and budget
movie_data <- movie_data %>%
mutate(profit = gross - budget,
return_on_investment_perc = (profit/budget)*100)

# Checking if movie color is an influential factor towards it's score

table(movie_data$color)

# It can be observed that the data in color is completely partial towards colored movies
# and thus it is not an influential factor and we can remove it
movie_data <- subset(movie_data, select = -c(color))

# Checking if color is removed from the data and returns a NULL value
movie_data$color

# Checking if movie color is an influential factor towards it's score

table(movie_data$language)

# It can be observed that the data in languages is completely partial towards english
  movies
# and thus it is not an influential factor and we can remove it

movie_data <- subset(movie_data, select = -c(language))

# Checking if language is removed from the data and returns a NULL value
movie_data$language

# Checking if the country the movie is produced in is an influential factor towards it's
  score

table(movie_data$country)

# Approximately 79% movies are from the US, 8% from UK and 13% from other
  countries
# Thus we collectively represent the movie locations as: US, UK, Others

levels(movie_data$country) <- c(levels(movie_data$country), "Others")

```

```
movie_data$country[(movie_data$country != 'USA') & (movie_data$country != 'UK')] <-
'Others'
movie_data$country <- factor(movie_data$country)
```

```
# Checking if only 3 locations are available
```

```
table(movie_data$country)
```

```
# DATA VISUALIZATION
```

```
# Histogram of movies released each year
ggplot(movie_data, aes(title_year)) +
  geom_bar() +
  labs(x = "Year movie release", y = "No. of movies released", title = "Histogram of no.
of movies released each year") +
  theme(plot.title = element_text(hjust = 0.5))
```

```
# It can be seen that the movie boom came after 1980 and thus we represent the data only
after 1980
```

```
movie_data <- movie_data[movie_data$title_year >= 1980,]
```

```
# Visualizing top 20 movies based on profits in Million$
install.packages("ggrepel")
library(ggrepel)
movie_data %>%
  filter(title_year %in% c(2000:2016)) %>%
  arrange(desc(profit)) %>%
  top_n(20, profit) %>%
  ggplot(aes(x=budget/1000000, y=profit/1000000)) +
  geom_point() +
  geom_smooth() +
  geom_text_repel(aes(label=movie_title)) +
  labs(x = "Budget in $million", y = "Profit in $million", title = "Top 10 Profitable
Movies") +
  theme(plot.title = element_text(hjust = 0.5))
```

```
# Using profits and return on investment variables are criteria to find 20 most profitable
movies
```

```
movie_data %>%
  filter(budget > 100000) %>%
  mutate(profit = gross - budget,
         return_on_investment_perc = (profit/budget)*100) %>%
  arrange(desc(profit)) %>%
  top_n(20, profit) %>%
  ggplot(aes(x=budget/1000000, y = return_on_investment_perc)) +
  geom_point(size = 2) +
  geom_smooth(size = 1) +
```

```

geom_text_repel(aes(label = movie_title), size = 3) +
xlab("Budget in $million") +
ylab("Percent of Return on Investment") +
ggtitle("20 Most Profitable Movies based on their Return on Investment")

# Visualizing 20 top directors based on the highest IMDB scores
install.packages("formattable")
library(formattable)
movie_data %>%
  group_by(director_name) %>%
  summarise(avg_imdb = mean(imdb_score)) %>%
  arrange(desc(avg_imdb)) %>%
  top_n(20, avg_imdb) %>%
  formattable(list(avg_imdb = color_bar("Red")), align = 'l')

# Plotting commercial success vs critical acclaim
movie_data %>%
  top_n(20, profit) %>%
  ggplot(aes(x = imdb_score, y = gross/10^6, size = profit/10^6, color = content_rating))
+
  geom_point() +
  geom_hline(aes(yintercept = 600)) +
  geom_vline(aes(xintercept = 7.75)) +
  geom_text_repel(aes(label = movie_title), size = 4) +
  xlab("IMDB Score") +
  ylab("Gross Money earned in million$") +
  ggtitle("Commercial Success Vs Critical Acclaim") +
  annotate("text", x = 8.5, y = 700, label = "High Ratings \n & High Gross") +
  theme(plot.title = element_text(hjust = 0.5))

# The above observation shows that there is hardly any correlation between critical
  acclaim and the movie's commercial success
# Visualizing relation between facebook likes and IMDB scores

library(plotly)
movie_data %>%
  plot_ly(x = ~movie_facebook_likes, y = ~imdb_score, color = ~content_rating, mode =
"markers",
  text = ~content_rating, alpha = 0.7, type = "scatter")

# Movies with high facebook likes can be seen to have higher IMDB score

# DATA PRE-PROCESSING

install.packages("data.table")
library(data.table)

```

```

# Find number of directors
sum(uniqueN(movie_data$director_name))

# Find number of actors
sum(uniqueN(movie_data[, c("actor_1_name", "actor_2_name", "actor_3_name")]))

# The names of the directors, actors 1 2 3 are so different that it will not contribute in
  predicting the score.
# The plot keyword is too diverse to be used as a predictor
# The movie IMDB link is redundant

movie_data <- subset(movie_data, select = -c(director_name, actor_2_name,
  actor_1_name,
  movie_title, actor_3_name, plot_keywords,
  movie_imdb_link))

# To avoid multicollinearity we remove the 2 previously added variables

movie_data <- subset(movie_data, select = -c(profit, return_on_investment_perc))

# Plot heatmap of the entire data as of now

ggcorr(movie_data, label = TRUE, label_round = 2, label_size = 3.5, size = 2, hjust =
  .85) +
  ggtitle("Correlation Heatmap") +
  theme(plot.title = element_text(hjust = 0.5))

# Based on the heatmap, we can see some high correlations (>0.7) between predictors.
# The highest correlation value observed is 0.95 and we can see that
  actor_1_facebook_likes is highly correlated with the cast_total_facebook_likes
# and both actor2 and actor3 are also correlated to the total.
# Thus we modify them into 2 variables: actor_1_facebook_likes and
  other_actors_facebook_likes.

movie_data$other_actors_facebook_likes <- movie_data$actor_2_facebook_likes +
  movie_data$actor_3_facebook_likes

# There is high correlations among num_voted_users, num_user_for_reviews and
  num_critic_for_reviews.
# We want to keep num_voted_users and take the ratio of num_user_for_reviews and
  num_critic_for_reviews.

movie_data$critic_review_ratio <- movie_data$num_critic_for_reviews /
  movie_data$num_user_for_reviews

# Delete Columns

```



```

movie_data <- subset(movie_data, select = -c(cast_total_facebook_likes,
actor_2_facebook_likes, actor_3_facebook_likes,
num_critic_for_reviews, num_user_for_reviews))

# Plotting heatmap to review post changes

ggcorr(movie_data, label = TRUE, label_round = 2, label_size = 4, size = 3, hjust = .85)
+
  ggtitle("Correlation Heatmap") +
  theme(plot.title = element_text(hjust = 0.5))

# No strong correlation of value greater than 0.7 observed

# The aim is to build a project wherein the model predicts whether the movie is good or
# bad. So bin the scores in four buckets: less than 4(Bad),
# 4-6(OK), 6-8(Good) and 8-10(Excellent)

movie_data$binned_score <- cut(movie_data$imdb_score, breaks = c(0,4,6,8,10))

# Rearranging the data and renaming the column to make it readable

movie_data <- movie_data[,c(9,4,5,14,12,2,3,13,1,6,10,7,8,11,15)]
colnames(movie_data) <- c("budget", "gross", "user_vote", "critic_review_ratio",
  "movie_fb", "director_fb", "actor1_fb", "other_actors_fb",
  "duration", "face_number", "year", "country", "content",
  "imdb_score", "binned_score")

# To apply models, splitting the data into training, validation and test sets with the ratio of
# 6:2:2
set.seed(45)
train.index <- sample(row.names(movie_data), dim(movie_data)[1]*0.6)
valid.index <- sample(setdiff(row.names(movie_data), train.index),
  dim(movie_data)[1]*0.2)
test.index <- setdiff(row.names(movie_data), union(train.index, valid.index))
train <- movie_data[train.index, ]
valid <- movie_data[valid.index, ]
test <- movie_data[test.index, ]

# IMPLEMENTATION OF ALGORITHMS

# CLASSIFICATION TREE

# Implementing a full grown tree

library(rpart)
library(rpart.plot)
# Full grown tree
class.tree <- rpart(binned_score ~ . -imdb_score, data = train, method = "class")

```

```

## plot tree
prp(class.tree, type = 1, extra = 1, under = TRUE, split.font = 2, varlen = 0)

# Implementing Best pruned tree

set.seed(51)
cv.ct <- rpart(binned_score ~ . -imdb_score, data = train, method = "class",
               cp = 0.00001, minsplit = 5, xval = 5)
printcp(cv.ct)
pruned.ct <- prune(cv.ct,
                  cp = cv.ct$cp.table[which.min(cv.ct$cp.table[, "xerror"]), "CP"])
length(pruned.ct$frame$var[pruned.ct$frame$var == "<leaf>"])
prp(pruned.ct, type = 1, extra = 1, split.font = 1, varlen = -10)

# Apply model on training set
tree.pred.train <- predict(pruned.ct, train, type = "class")
# Generate confusion matrix for training data
confusionMatrix(tree.pred.train, train$binned_score)

# Apply model on validation set
tree.pred.valid <- predict(pruned.ct, valid, type = "class")
# Generate confusion matrix for validation data
confusionMatrix(tree.pred.valid, valid$binned_score)

# Apply model on test set
tree.pred.test <- predict(pruned.ct, test, type = "class")
# Generate confusion matrix for test data
confusionMatrix(tree.pred.test, test$binned_score)

# K NEAREST NEIGHBOUR

library(FNN)
# Using model.matrix() to create dummy variables for country and content
movie_data2 <- movie_data
movie_data2$country <- as.factor(movie_data2$country)
movie_data2$content <- as.factor(movie_data2$content)
movie_data2[,c("country_UK", "country_USA", "country_Others")] <- model.matrix(~
  country - 1, data = movie_data2)
movie_data2[,c("content_G", "content_NC17", "content_PG", "content_PG13",
  "content_R")] <- model.matrix(~ content - 1, data = movie_data2)

# Select useful variables for future prediction
movie_data2 <- movie_data2[, c(1,2,3,4,5,6,7,8,9,10,11,16,17,18,19,20,21,22,23,15)]
# Partition the data into training and validation sets
set.seed(52)
train2 <- movie_data2[train.index, ]
valid2 <- movie_data2[valid.index, ]

```

```

test2 <- movie_data2[test.index, ]

# Initializing normalized training, validation, test data, complete data frames to originals
train2.norm <- train2
valid2.norm <- valid2
test2.norm <- test2
movie_data2.norm <- movie_data2

# Using preProcess() from the caret package to normalize predictors
norm.values <- preProcess(train2[, -20], method=c("center", "scale"))
train2.norm[, -20] <- predict(norm.values, train2[, -20])
valid2.norm[, -20] <- predict(norm.values, valid2[, -20])
test2.norm[, -20] <- predict(norm.values, test2[, -20])
movie_data2.norm[, -20] <- predict(norm.values, movie_data2[, -20])

# Finding the best K
# Initialize a data frame with two columns: k, and accuracy

accuracy.df <- data.frame(k = seq(1, 20, 1), accuracy = rep(0, 20))

# Computing knn for different k on validation data

for(i in 1:20) {
  knn.pred <- knn(train2.norm[, -20], valid2.norm[, -20],
                  cl = train2.norm[, 20], k = i)
  accuracy.df[i, 2] <- confusionMatrix(knn.pred, valid2.norm[, 20])$overall[1]
}
accuracy.df

# Applying model on test set

knn.pred.test <- knn(train2.norm[, -20], test2.norm[, -20],
                    cl = train2.norm[, 20], k = 9)

# Generating confusion matrix for test data

accuracy <- confusionMatrix(knn.pred.test, test2.norm[, 20])$overall[1]
accuracy

# RANDOM FOREST

install.packages("randomForest")
library(randomForest)
set.seed(53)
rf <- randomForest(binned_score ~ . -imdb_score, data = train, mtry = 5)
# Show model error
plot(rf)
legend('topright', colnames(rf$err.rate), col=1:5, fill=1:5)

```

```

# Get importance
importance <- importance(rf)
varImportance <- data.frame(Variables = row.names(importance),
                             Importance = round(importance[, 'MeanDecreaseGini'], 2))

# Creating a rank variable based on importance
rankImportance <- varImportance %>%
  mutate(Rank = paste0('#', dense_rank(desc(Importance))))

# Using ggplot2 to visualize the relative importance of variables
ggplot(rankImportance, aes(x = reorder(Variables, Importance),
                           y = Importance, fill = Importance)) +
  geom_bar(stat = 'identity') +
  geom_text(aes(x = Variables, y = 0.5, label = Rank),
            hjust = 0, vjust = 0.55, size = 4, colour = 'red') +
  labs(x = 'Variables') +
  coord_flip() +
  theme_few()

install.packages("caret")
library(caret)
set.seed(632)
# apply model on validation set
rf.pred.valid <- predict(rf, valid)
# generate confusion matrix for validation data
confusionMatrix(rf.pred.valid, valid$bin_score)

```